

Project Alterac Valley : Technical Design Document (Gameplay)

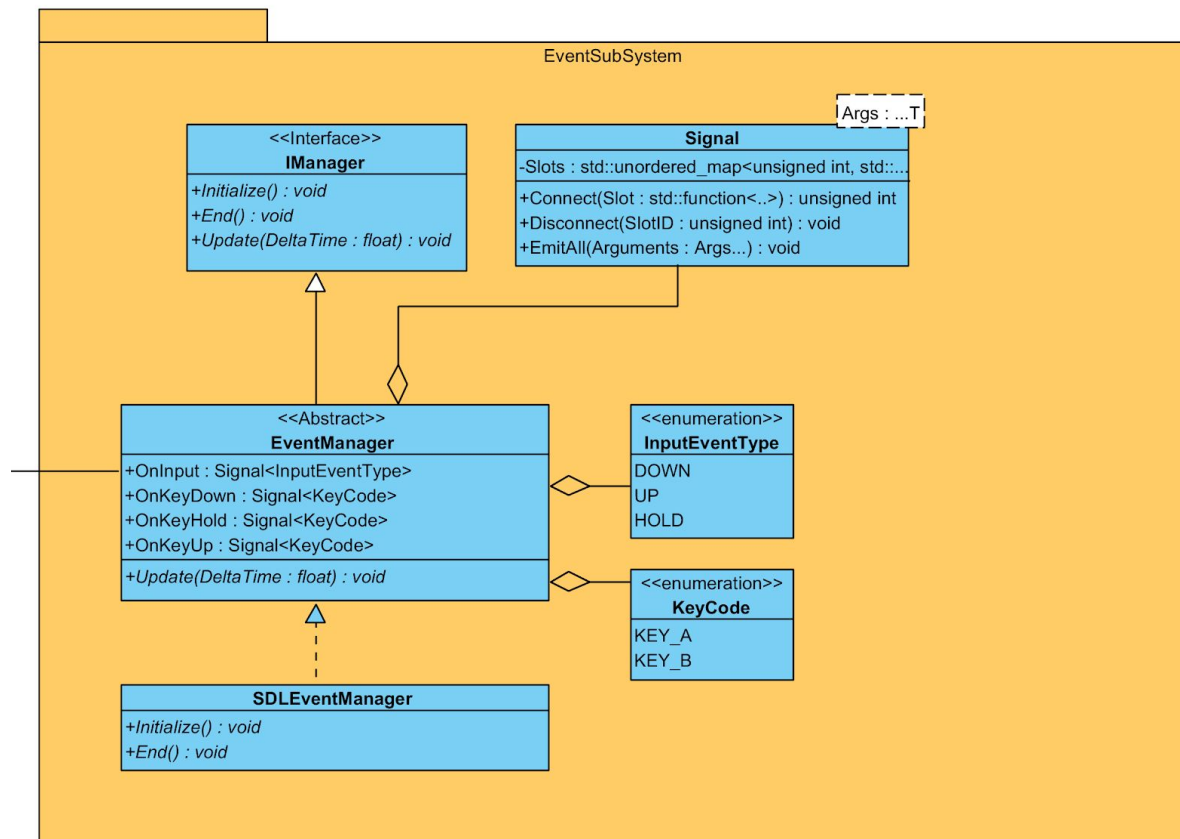
Landon Grant
Jean-François Carnovale
Jaymie Xu
Brandon MacLean
Nathan Ellenberger

General Architecture	3
Game Entry Logic	3
System of Entities and Components	4
Scene Management	5
Global Game Scripts	6
Unit Class and Player Logic	7
Unit Logic Component	7
Unit Status System	8
Unit Behaviour / State Machine	9
Collision and Physics Components	10
Event Handle And Data Manager	11
Input/Event Handle	11
Data Manager and Server Communication	12
Rule Validation	14
UIManager	14
Scene Information Splitting	15
Diagram & Game Loop explained	17
Server	18
Main Game Loop	18

General Architecture

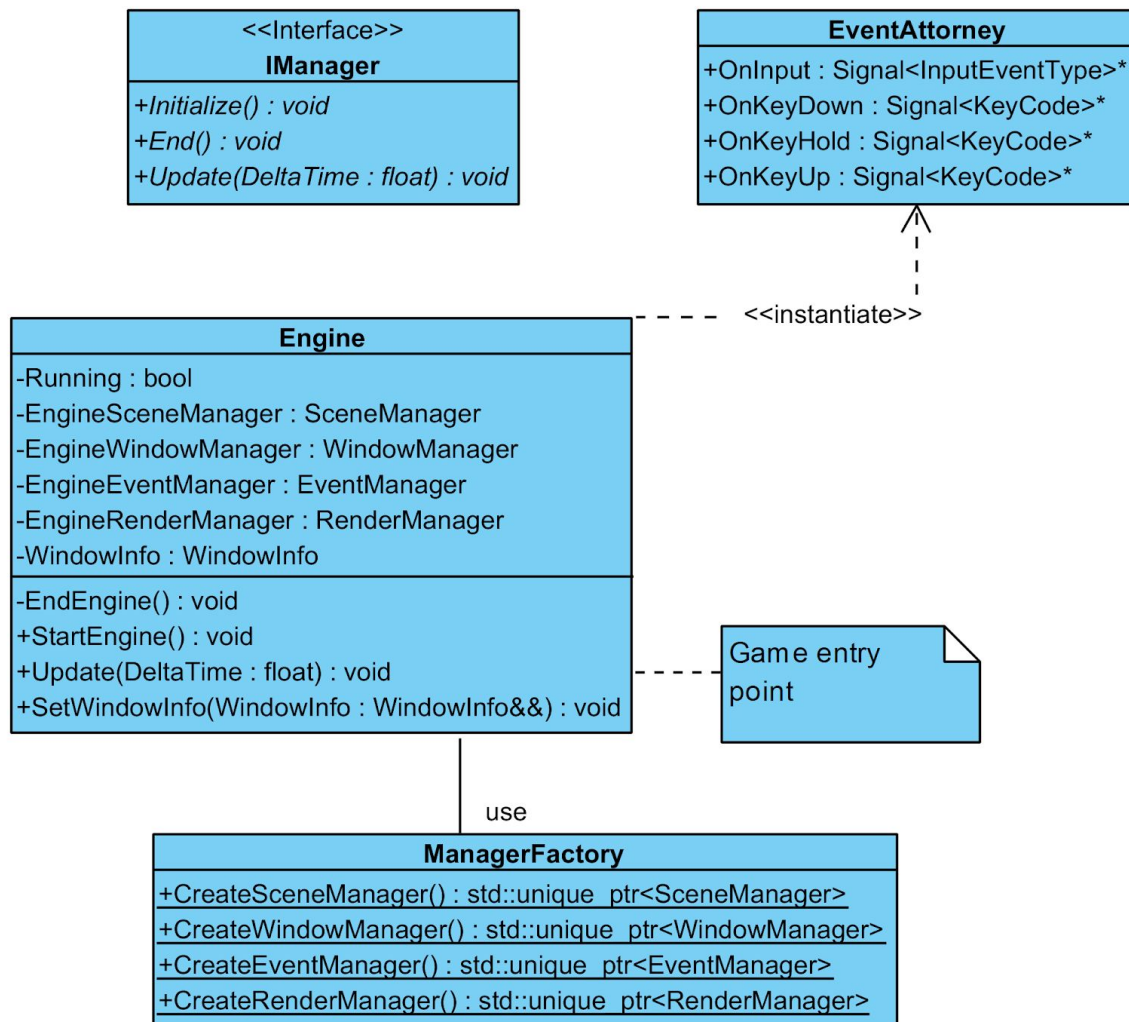
Game Entry Logic

Project Alterac Valley will be coded in C++. At this stage of development, most teams have agreed upon using OpenGL and SDL as two of our main graphic and event handling API. Dependencies such as OpenGL and or SDL may change later on (For argument sake, we decided to use DirectX and GLFW instead in future), therefore we decided to heavily rely on the dependency inversion principle for most of our engine subsystems.



(Full UML exported as zoomable .pdf files)

As an example, The EventManager class is an abstract class with some common functionalities. In order to handle event a programmer has to use the EventManager as a base class and write their own event manager; in this case we wrote SDLEventManager since SDL is the event api dependency here. During the actual creation of the engine there will be another class called ManagerFactory which will construct the appropriate concrete EvenManager base on precompile macro settings.



(Full UML exported as zoomable .pdf files)

System of Entities and Components

We decided to make Project Alterac Valley's inner working system a heavily modular one. This means that everything, aside from some of the major managers and engine code, should be either a **GameObject** (aka an entity) or a **Component**. To illustrate this, a playable character is nothing more than just a **GameObject** with **MeshRendererComponent**, **AnimationPlayerComponent**, **UnitComponent**, **UnitStatComponent**, and finally a **CollisionComponent**. On the other hand an AI player may exactly the same as a playable character, with one exception: It also has a **AIControllerComponent**. Without this system of

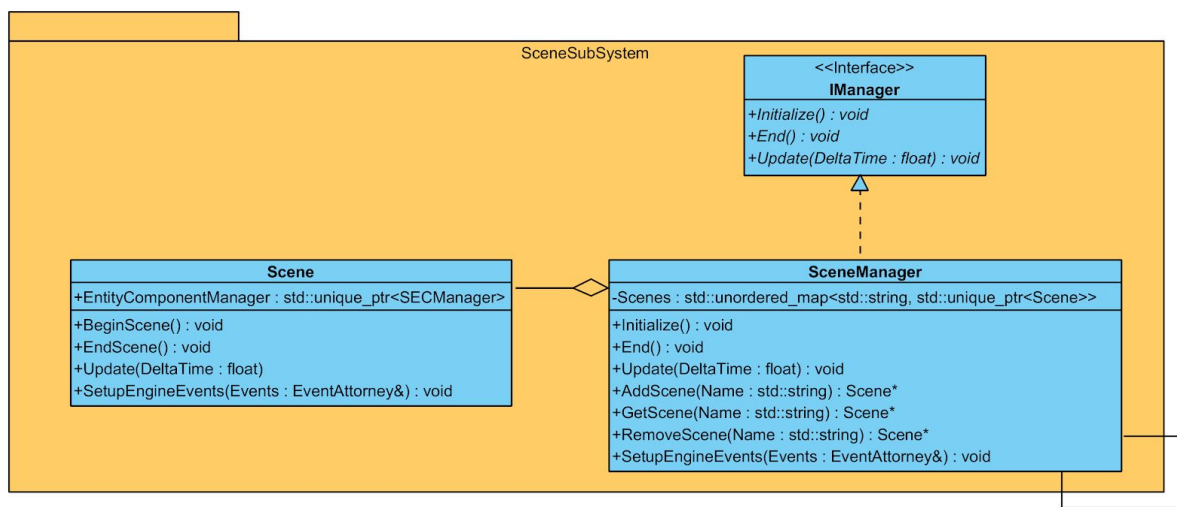
entity and components such behaviour can only be achieved with complex strategy pattern or with heavy usage of inheritance and repeated code.

Each IComponent object will use a GUID to identify their type, for example a RenderComponent's GUID is 1 and a CollisionComponent GUID is 2. To generate this ID we used CRTP(Curiously Recurring Template Pattern) to allow derived class to pass their type by template into super class. Once the is in the super class, the system would then call a static GUID class to generate an ID based on an object's type(See UML for detail). We would employ several template meta programming tricks as well as local static variable to achieve this effect.

In order to make sure our components are store continuously in memory. We would be creating all component by ComponentManager which will store then in a hashmap of vectors the key being the component GUID. GameObjects will have pointers to all their components thus creating the illusion that GameObjects Owns Component(GameObject can easily access their own component) while in truth ComponentManager manages it all. This also allows search by component type easily implemented(Using the hashmap get Component type GUID to fetch the correct vector of components).

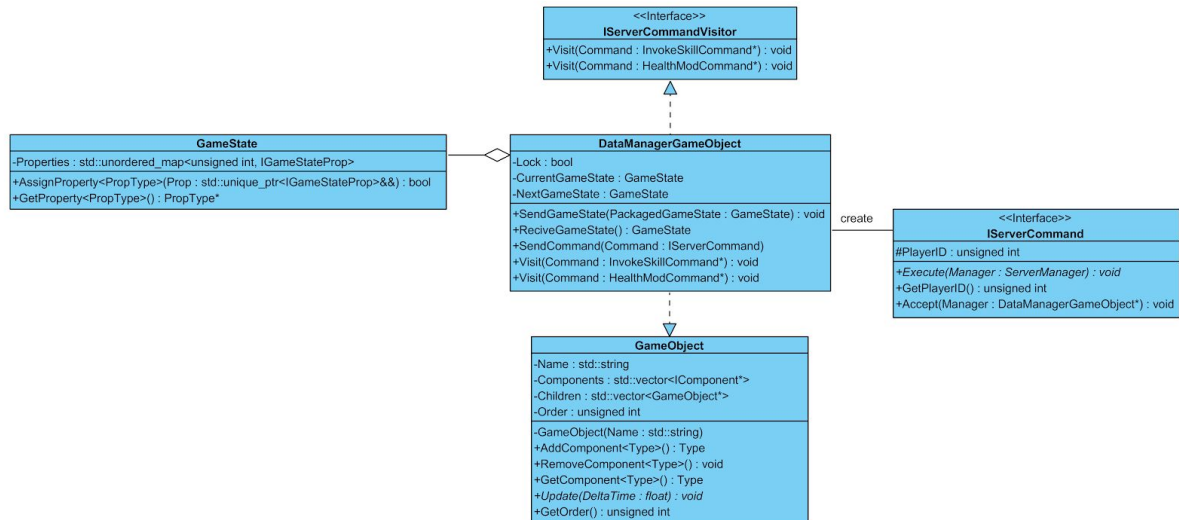
Scene Management

Unlike some of our previous works. Project alterac valley's Scene class will not be an inheritable abstract class(This decision may likely change in future depending on people's ability to comprehend c++ RTTI). Making scenes inheritable allows direct modification inside them but all the c++ code inside a inherited scene is hard coded and would be difficult to change certain thing at runtime. Hence we decided to make the scene a container used to update all GameObject and Components it owns. To actually put GameObjects into a scene you would use editor to design a level based on a list of Components and GameObjects. And during the runtime we would use RTTI to determine what GameObject and Componets a scene needs, keep in mind that this is only invoking RTTI once during a scene startup so the performance hit would only be some extra loading time, however doing it this way adds a lot more flexibility. *We do recognize that RTTI is a lot to take in so we designed our system to allow the inheritance of the scene class in future if needed.*



Global Game Scripts

Because of the way entity, components, and scenes work all level game managers(These are different than engine managers) will need to be a custom gameobject. For example, DataManager is called DataManagerGameObject because so that it allows the injection of the server update logic into scene.



(Full UML exported as zoomable .pdf files)

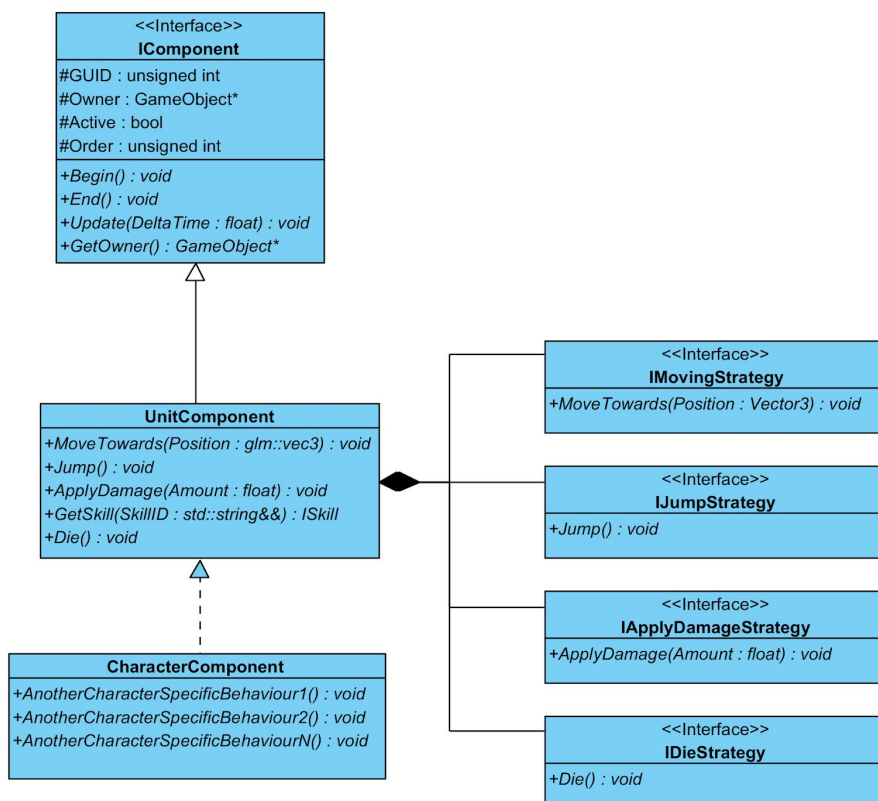
Unit Class and Player Logic

Unit Logic Component

All the characters (Players and NPCs) controlling logic in the game will use a class (or one of its specialized subclass) called UnitComponent (Inherited from IComponent interface; everything is a component or Gameobject). The UnitComponent base class will have common behaviour functions such as Moving, Dying, AutoAttack and Respawn. UnitComponent will also have some basic members such as a mesh, animations etc. Since every player class and NPC will have some similarities and some differences we decided to use the strategy pattern to achieve this effect.

For example, a Knight class can be made of a UnitComponent class with MoveStrategy1, DieStrategy2, AutoAttackStrategy2, and RespawnStrategy4 while a Mage class can be made of a UnitComponent class with MoveStrategy1, DieStrategy3, AutoAttackStrategy2, RespawnStrategy1. The Knight and Mage both use the same MoveStrategy and AutoAttackStrategy by using the strategy pattern we essentially avoid the need of repeating unnecessary code. Furthermore, because most Class/NPC behaviours are wrapped inside Strategy classes we can use a factory class to build the desired class at runtime.

This flexibility enables us the ability to easily define class archetype in a JSON file (see example below). Should there ever be a need to change these classes' behaviours in the future, all we need to do is modify the JSON file and change a single line of our game code.

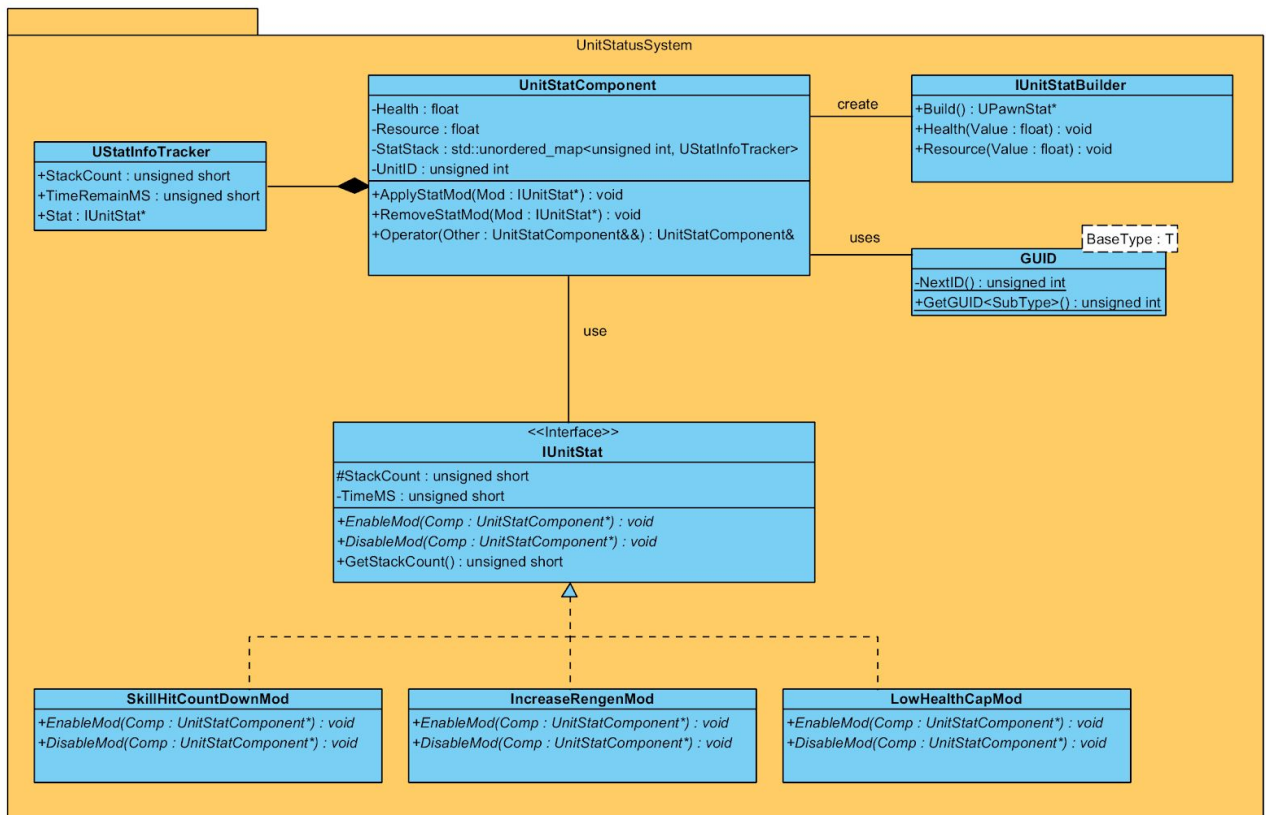


We also recognized that there will be some cases where defining specific traits of a certain archetype may be difficult when using the strategy pattern above. As such, all the base functions in the character class like Move, AutoAttack, Die, and Respawn can be overwritten. This allows us to derive a subclass from the UnitComponent class, for example: UnitComponent (Derive)-> SpecialUnitComponent. And since UnitComponent is a child class from the base class UnitComponent, it may also use the same strategy classes while also having the ability to override or add functionalities.

Lastly, skills are handled differently from strategies. Each UnitComponent object will contain a list (std::vector object) of Skill based object. A Skill object contains some basic information about a skill/ability such as if the damage range is area of effect or single target. Skill class will also have a mandatory pure virtual function to define how a skill should be invoked (does it shoots a projectile or does it spawn a bunch of fireball that rain down from sky etc).

Unit Status System

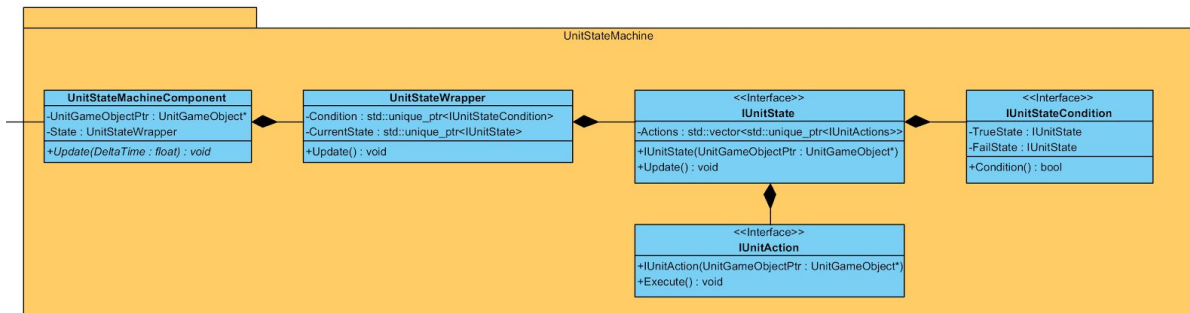
Every single Unit be it player or AI can have some kind of status on them at any given point. A player may have a fast regen stat and a quick cast stat, it may be tempting to create 2 sets of variable for each status: current_health, max_health, current_cast_mod, base_cast_mod, etc. This approach would work for small projects; however, we are dealing with an MMO and there are all sorts of buffs, even skill from another player is a “debuff” which by the end of that buff would case it's attached player to lose some health. Therefore its is crucial that we make the buff system as modular as possible. Instead of storing two variable for each stat a player may have we created a component called UnitStatComponent and a base interface class called IUnitStat, every buff/debuff including health regen will be derived from IUnitStat and this interface has 2 primary functions EnableMod and DisableMod both takes in a parameter of UnitStatComponent and have their own rules of modifying it. Whenever there is a new buff UnitStatComponent calls ApplyStatMod and when ever that buff is over UnitStatComponent can call RemoveStatMod. That way UnitStatComponent is decoupled from all the possible buff that could ever be done to a character. Further more this also allows buff stacking as you can call ApplyStatMod on the same buff multiple times.



(Full UML exported as zoomable .pdf files)

Unit Behaviour / State Machine

On the other hand, a unit has to behave a certain way. By that we mean a player should be able walk run when it is currently not stunned, similarly an AI should not mindlessly wonder randomly on map then eventually gets kill by player. The easiest way to solve this problem is to use a state machine. Each character player/ai would have a specific dynamically set states. For example player start off in moveable state which can transition into a stunned state or a casting state; all state may transition to a death state. This way, with server validation system, player can not move while getting stunned unless the server tells the player to transition back from stunned state to moving state. In contrast the AI may use the statemachine in a different manner, instead of telling the AI what they can do, the AI state machine would dictate what they could do. For instance an AI begins with patrol a certain area and will only transition to chase state(sees player) if it encounters a player.

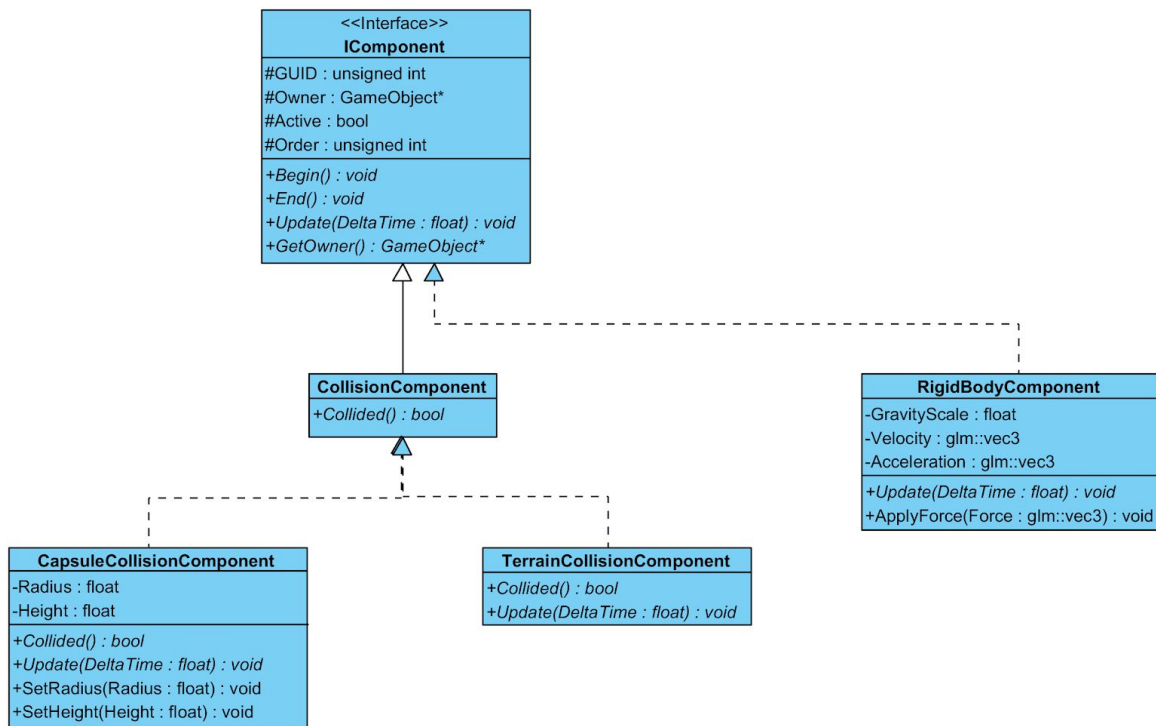


(Full UML exported as zoomable .pdf files)


Collision and Physics Components

Component classes are all derived from the super interface IComponent. The amount of component are expected to be increase as the development proceeds. At this moment the primary Components used by Gameplay, aside from all the unit components, are the collision and physics components.

First, all collision components will inherit from an abstract call called CollisionComponent which would provide an overridable function called Collided which will check if this compoent has detected a collision with another collider component of any type and return a boolean value. So far we have CapsuleCollisionComponent which takes in a height and a radius, this will most likely use its bottom sphere to check collision with another object. The other collider we have was the terrain collider which would detect if any object has come into contact with the terrain.



There are many methods used in games to detect collision. Examples of such are GJK SAT AABB, since MMO is not a physics heavy game we decided to go for the least costing algorithm AABB collision detection for our terrain. AABB stands for Axis Aligned Bounding Box, it utilizes 2 bounding boxes generated from 2 object to detect if they ever need to be collided. Character are most likely to be collided with terrain, however they utilizes a capsule collider(This is used to get a somewhat smooth motion while going up a ramp) AABB cannot directly act upon a capsule and a bounding box/volume(We need to have 3 axis overlap). Therefore we need to first take the bottom most sphere of the capsule approximate a bounding volume if AABB passes we uses the capsule collider to further adjust the character's position, otherwise the character and the terrain are not collided.



Lastly the physics component in our case(Again because its an MMO) would have one job: pulling the player down due to gravity.

Event Handle And Data Manager

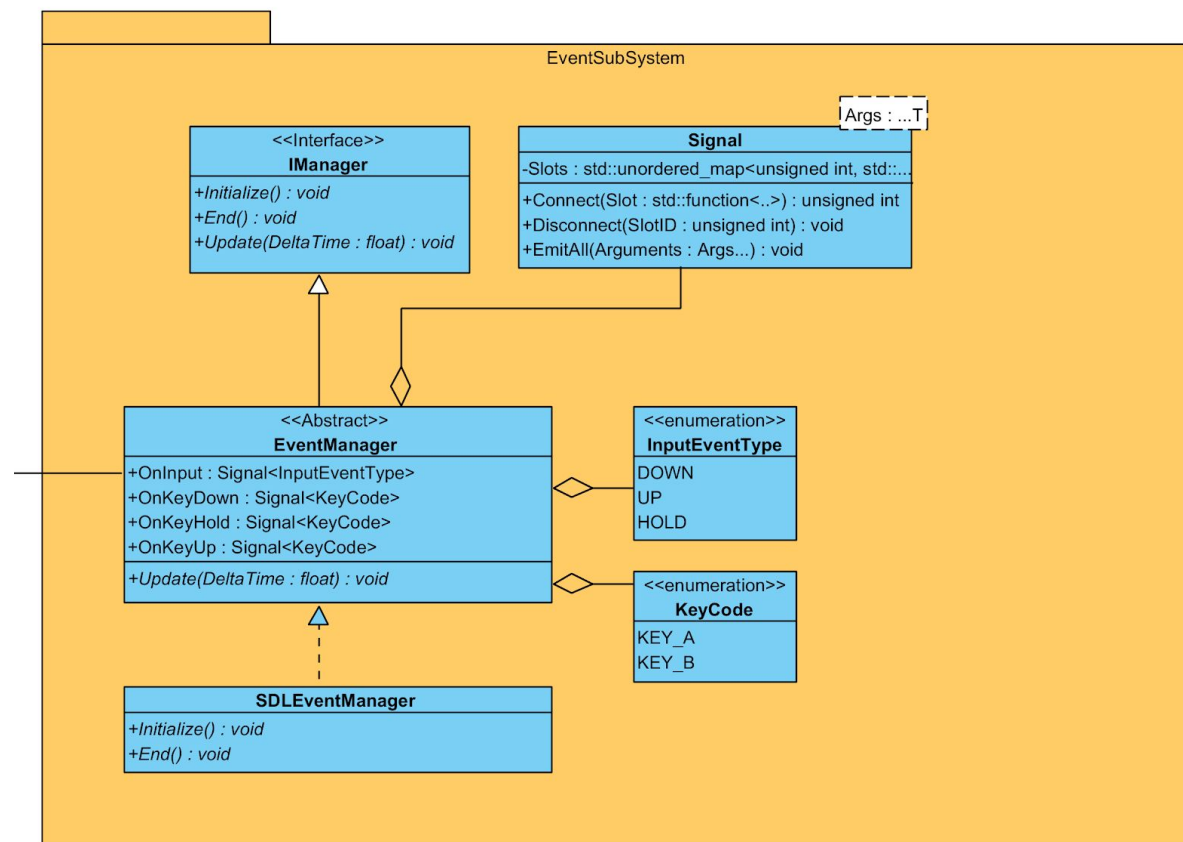
Input/Event Handle

Events will be handled by EventManager and as mentioned previously, it is an engine level manager. At this moment we have decided to use SDL as our event handling API, to make sure that SDL do not couple with our own code we have safely wrapped it inside a EvenManager subclass called SDLEventManager. This way all the engine needs to know is a pointer to EvenManager. The actual event is forwarded by a system called Signals and Slots(This was used in QT: a GUI lib for c++). Signals and Slots are very much like C#'s delegate and a observer pattern that which does not required the user to set observer class as a super class. A signal is essentially a function pointer, slots are all the function registered to it for now we have 4 signals OnInput, OnKeyDown, OnKeyHold, OnKeyUp, of course more can be added later such as mouse input. In order to avoid singleton and allow event manager to leak into every part of our engine, we decided to use dependency injection. First we made an EventAttorney class that hides some of the EventManager functions such as Update(User of the engine should not have the ability to force called Update function that was handled by the Engine class, yes people have tried to do this before to my engine)

EventAttorney also has pointer to all the Signal object the EventManager class has so essentially the EventAttorney class is like a proxy class - a simplified version of a complex object.

To actually handle event in a Game level one must called the function SetupEngineEvents(EventAttorney& Events) on GameObject or Component, the EventAttorney here is the injected dependency and you would now need to bind a function to one of the signal. For example
Events.OnKeyDown.Connect(&Object::HandleKeyDown(KeyCode));

Object::HandleKeyDown(KeyCode e) { if (e == KEY_A) { ... } }



(Full UML exported as zoomable .pdf files)

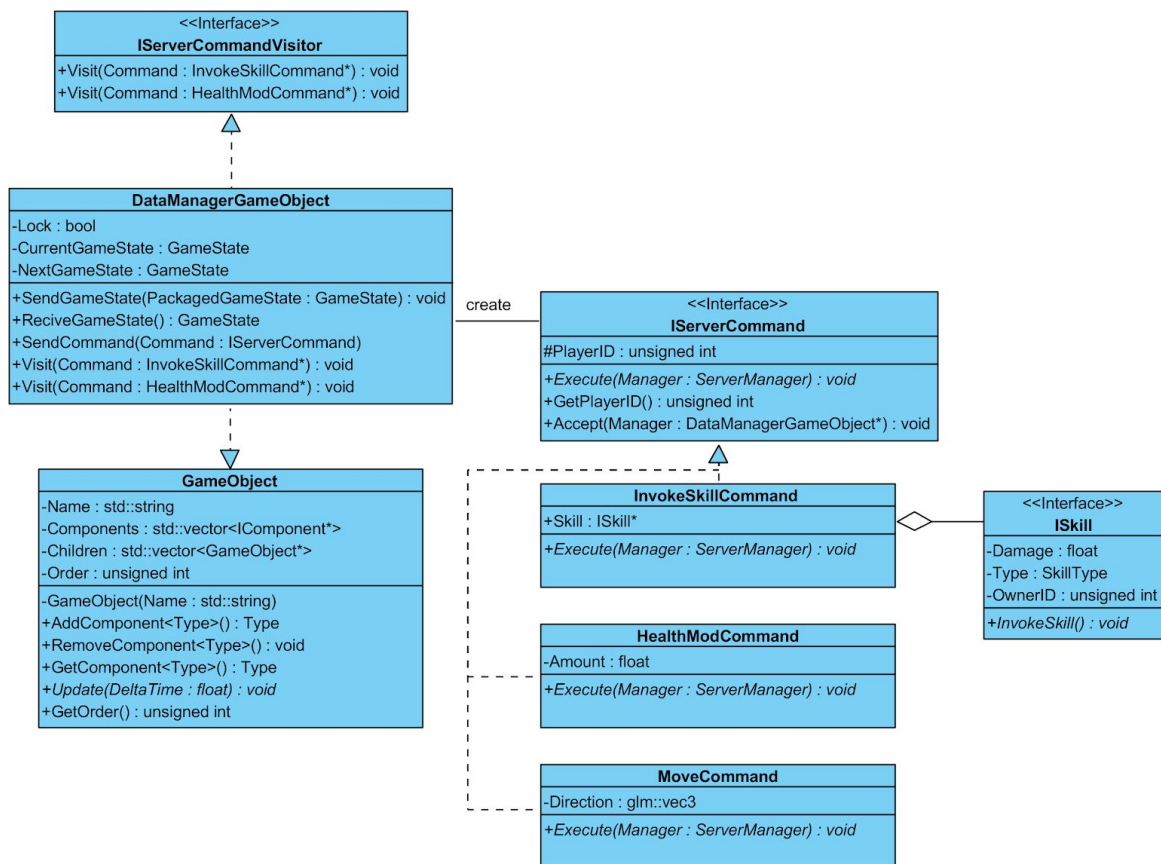
Data Manager and Server Communication

All the communication between the client and the server will be handled by a class called **DataManagerGameObject** (This is not an engine level manager thus it has to be a **GameObject** at least). **DataManagerGameObject** would send out one instance of the current game state of the current player (Health, Position, State Machine State) and a series of **IServerCommand** objects. A player may not directly ask to change their current state, if a player wishes to regain health, the client side must send out a **IServerCommand** based object like **HealthModCommand** and request the server to change this state for them. On the other hand, the server will implement a visitor pattern to determine which command the

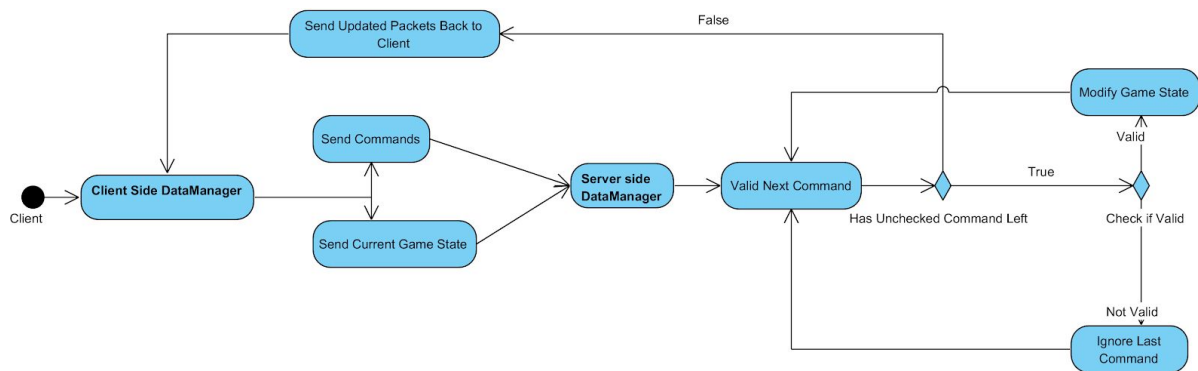
player just send and check whether or not this command is currently valid(See below for rule validation). If not valid server does nothing, if the command is valid the server would call `Command.Execute(this)`; to modify the state. After the server finishing executing all commands for a player it returns the newly modified `GameState` back to player.

The `DataManager`, along with other network related objects, would most likely operate on a different thread than the main thread. Therefore it must be thread safe, whenever `DataManager` is Sending datas to the servers we will switch the current `Game state` pointer so client does not accidentally write to a `GameState` that is currently transfer to the server.

One important thing to note here is that this communication does not happen every frame. Therefore, on the player's end, if a server didn't respond they may still able to move around before the server kicks back in and rubber bands their position and stats back.




(Full UML exported as zoomable .pdf files)



Rule Validation

This static class is used to validate player position and inputs made by the player. For example, if the player is stunned, rather than handling inputs normally, it should prevent the associated action from taking place unless appropriate and instead display a message and play a sound effect notifying the player that they cannot make that action due to being stunned.

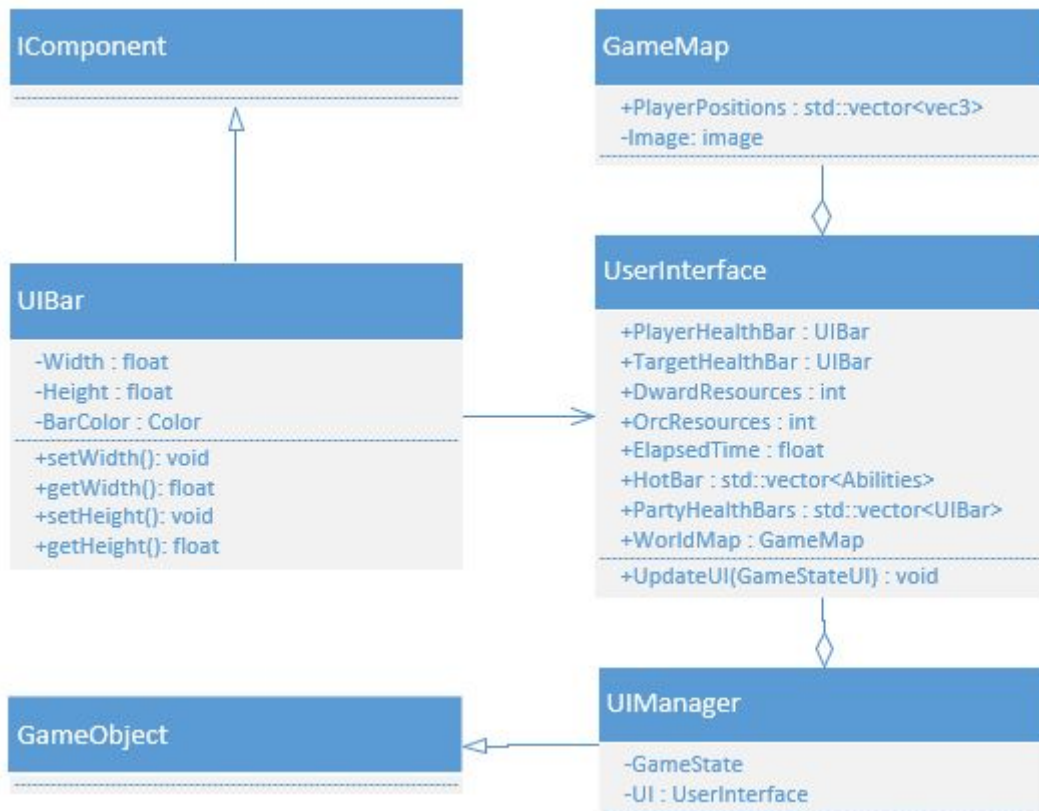


The position validation will prevent the client from getting out of sync with the server as well as prevent cheating. The class has a MaximumVariantDistance threshold which, if exceeded, will cause the player character on the client to rubberband back to its position on the server. It works by taking the time elapsed since the last update and checking how much the unit has moved since then. Depending on that units abilities, it may be able to travel a large amount of distance in a short time frame. This is where the variant distance comes into play, allowing a bit of leeway before deciding to determine the movement as invalid.

UIManager

The UIManager gets the information of the GameState and breaks up the information to be appropriate to Update the UI, stores it into a GameStateUI and sends it to the UserInterface by calling the public UpdateUI method on the UserInterface. The UserInterface handles this

information and updates each of its members as needed. For example, the width and value of the HealthBar will be adjusted to match that of the GameState.



Scene Information Splitting

As the game will have various Game objects of different types that will all need to be persistent and accounted for at all time, that inherently means that there will be a lot of information in the current game scene at every iteration. With such a large amount of information, not all of it would necessarily be relevant to every type of entity in the scene, having each entity in the scene receive information from every other entity and having to do whatever type of comparison operations with those entities is unrealistic, and can cause severe performance issue. However, there are ways to split up that information in a way where only the most relevant information of the various game objects relative to each objects current location will be checked. The best way to achieve this goal is to store references to the game objects based on their position in the world space inside of a binary space partitioning tree (BSP tree). There are 3 main BSP tree variations that we looked at as a solution to this particular problem.

There are octrees, which are trees where each node has either no children and therefore stores the information pertaining to a set amount of game objects within its space, or it has a node that represents the center of its three-dimensional space and exactly eight children octrees that represent the eight quadrants that divide the three-dimensional space represented by the parent octree. The advantage to the octree is most relevant in cases

where there is important information that can be found in three-dimensional space and you would like to use every entity that resulted from the search parameter.

There are k-d trees, which are k dimensional trees. These trees function in a way where it is divided by dimensional axis and cycles through the dimensions per partition. The location of the partition is also not uniform as in order to make a well balanced k-d tree you would divide the partition at the median of all the data points that are used for that section of the tree. This makes k-d trees well balance and quick as the information stored in the tree will be of similar distance to the root node. However, as you need the median of the information stored in the tree, you would need to give it the information upfront and therefore variable information and addition of new nodes will result in an imbalanced tree and reduce its performance. The advantage to a k-d tree comes out mostly in the case of things like collision detection with persistent objects as their locations are pre-defined. Another advantage is that the way information is partitioned in a k-d tree is by a greater than, less than relationship as any node that is less than the median will be on one side of the partition where things that are greater will be on the other, as such in cases like collision or render culling these trees are advantageous, as they can easily discard anything that is at a greater distance and as such of no concern.

The final type is the quadtree, quad trees are typically used in two-dimensional space and are the two-dimensional equivalent of the octree. A quadtree has either a node that stores a set amount of information and no children or contains a reference to exactly four sub children. Each quadrant contains either its own set amount of information or its own reference to four of its own children. The advantages are similar to the octree as well however applied in two-dimensional space, this also means that they aren't as large as the octrees if you do not require three-dimensions of information. They can be used to easily determine whether or not something is outside of a relevant distance to a point or area that you are searching for information in and can also be used for collision and render culling.

For the sake of this game we have decided that a quadtree would be our best strategy in terms of performance given the context of our game. While the game may be a three-dimensional game, and therefore the initial choice would be to use octrees, the way our game is played, the information in the third dimension isn't very relevant in terms of deciding which objects need to be rendered or which objects have information that is relevant to the current object querying, as there is very little three-dimensional movement in the game and most things will be relatively close in the z-axis. As such the quad tree is smaller and faster than a octree and will serve the same purpose to reduce the number of operations done between each update of the game cycle. Given the proper ranges we will be able to easily divide the information in the game scene into multiple quadtrees and properly cull any information that is too far to be of concern to units.

The information for these trees will primarily be stored on the server as it will contain all of the objects that are in the game scene. At each update the tree will be recreated as with any movement the tree will need to be rebuilt in order to update all the positions. As for the passing of the information, players will query the server for any relevant information to them and their game state, the server will receive the players location, use it to generate a area that will be used to query the quad tree for a list of every unit that falls within that range, and then send all that information to the client. The client will then compare the list of units and update their game state accordingly so as to only include the units that are within that range and their updates at that pass, in order to save performance and cut back on the amount of information that is being passed between each object at each game state update. Thus we will have one large tree that will contain all the information on the units in the game, that only the server will have access too, and will query search radiuses around each unit in order to return to each unit only the relevant information to them and their game state.

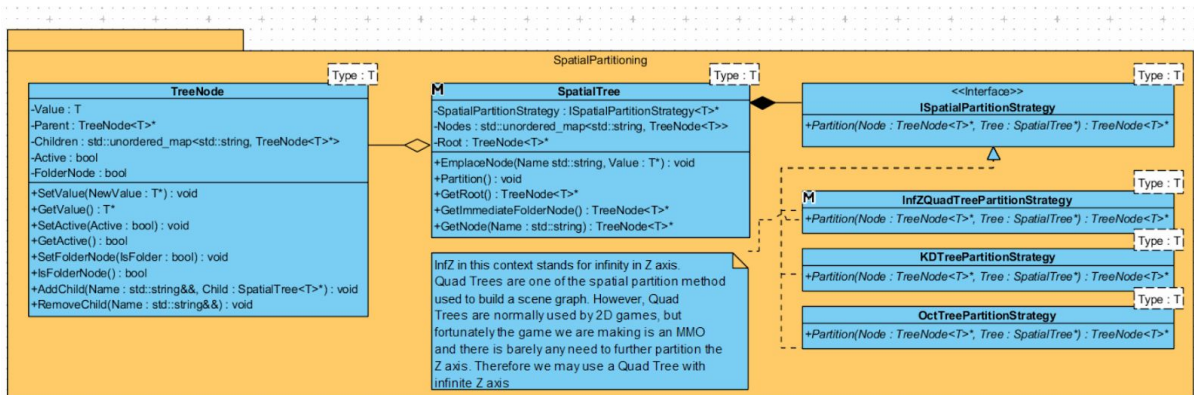

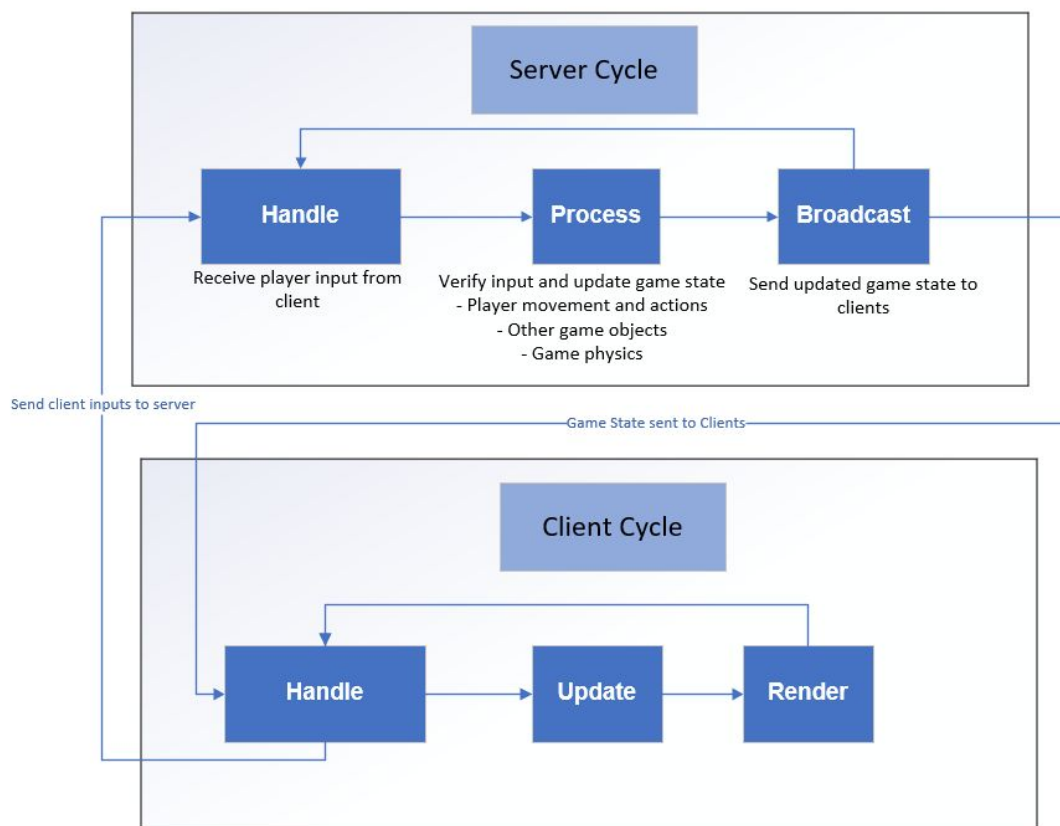


Diagram & Game Loop explained





Server

The server is going to be set up and everyone will be able to connect to it.

Broadcast/Handle: The server will first receive input/actions the player is doing which will be sent to the server via the main game loop. The server will take this input and then process it and check whether the move that was performed is valid and then let the player perform the action. The server will combine all certain packets of input received and validate it so that the next update call everyone can be in sync and the server has no problems.

Main Game Loop

Update(): The update function will constantly be changing the players position, status, etc. It will constantly be updating the game state to keep the player aware of everybody's actions being performed. If the player is out of sync with the server at some point in the update call, it will sync the player with the rest of the game.

Every time the game updates, the four main managers have their update functions called. SceneManager's Update will cause the the components of game objects within the scene to update, while EventManager's Update will update the event that the EventAttorney class passes to components when checking for events. WindowManager will update the game window in case any changes are made to it, and RenderManager will make any necessary changes to renderables in the game.

Handle Events(): This method will be constantly checking every loop if there was a player input and then it will carry out the appropriate action associated with the key press if all goes correctly. These events will be sent off to the server to be validated. Our handle events function is similar to the way that unreal engine handle's inputs. There is going to be a function within the component class called SetupEngineEvents which will be using dependency injection to bring in an EventAttorney class that contains the details of the event that occurred. E.G. EventAttorney.onKeyboardinput.register(your function here).

Render(): Our render function will be incorporated inside the main game loop and will be handling all the rendering of our game scene.

